

# UNIX System Programming

## Files and Directories

1

## File Pointer

- Both read() and write() will change the file pointer.
- The pointer will be incremented by exactly the number of bytes read or written.

2

## lseek

```
#include <sys/types.h>
#include <unistd.h>
off_t lseek( int fd, off_t offset, int whence );
```

- Repositions the offset of the file descriptor *fd* to the argument offset.
- *whence*
  - SEEK\_SET
    - The offset is set to offset bytes.
  - SEEK\_CUR
    - The offset is set to its current location plus offset bytes.
  - SEEK\_END
    - The offset is set to the size of the file plus offset bytes.


3

## lseek: Examples

- Random access
  - Jump to any byte in a file
- Move to byte #16
  - `newpos = lseek( file_descriptor, 16, SEEK_SET );`
- Move forward 4 bytes
  - `newpos = lseek( file_descriptor, 4, SEEK_CUR );`
- Move to 8 bytes from the end
  - `newpos = lseek( file_descriptor, -8, SEEK_END );`


4

## lseek - SEEK\_SET (10)



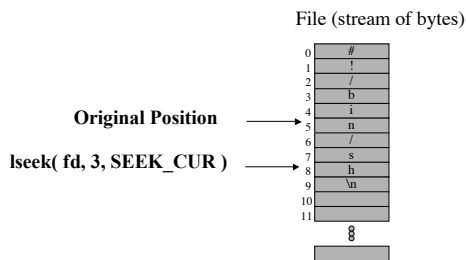
5

## lseek - SEEK\_CUR (-5)



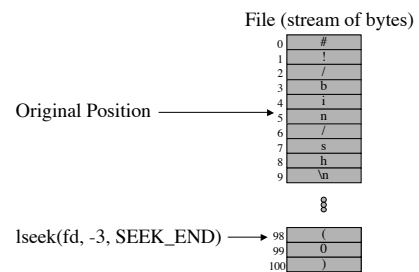
6

## lseek - SEEK\_CUR(3)




7

## lseek - SEEK\_END (-3)




8

## Read – File Pointer



9

## Write – File Pointer



10

## Example #1: lseek

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

char buf1[] = "abcdefghij";
char buf2[] = "ABCDEFGHIJ";

int main(void)
{
    int fd;

    if( (fd = creat("file.hole", 0640)) < 0 )
    {
        perror("creat error");
        exit(1);
    }
}
```

11

## Example #1: lseek (2)

```
if( write( fd, buf1, 10 ) != 10 )
{
    perror("buf1 write error");
    exit(1);
}

/* offset now = 10 */
if( lseek( fd, 40, SEEK_SET ) == -1 )
{
    perror("lseek error");
    exit(1);
}

/* offset now = 40 */
if( write( fd, buf2, 10 ) != 10 )
{
    perror("buf2 write error");
    exit(1);
}

/* offset now = 50 */
exit(0);
}
```

12

## File control of open files: fcntl()

```
#include <unistd.h>
#include <fcntl.h>

int fcntl( int fd, int cmd );
int fcntl( int fd, int cmd, long arg );
int fcntl( int fd, int cmd, struct lock *ldata )
```

- Performs operations pertaining to *fd*, the file descriptor
- Specific operation depends on *cmd*

13

## fcntl: cmd

- F\_GETFL
  - Returns the current file status flags as set by open().
  - Access mode can be extracted from AND'ing the return value
    - return\_value & O\_ACCMODE
    - e.g. O\_WRONLY
- F\_SETFL
  - Sets the file status flags associated with fd.
  - Only O\_APPEND, O\_NONBLOCK and O\_ASYNC may be set.
  - Other flags are unaffected

14

## Example 1: fcntl()

```
#include <stdio.h>
#include <sys/types.h>
#include <fcntl.h>

int main( int argc, char *argv[] )
{
    int accmode, val;

    if( argc != 2 )
    {
        fprintf( stderr, "usage: a.out <descriptor#>" );
        exit(1);
    }

    if( (val = fcntl(atoi(argv[1]), F_GETFL, 0)) < 0 )
    {
        perror( "fcntl error for fd" );
        exit( 1 );
    }

    accmode = val & O_ACCMODE;
```

15

```
if( accmode == O_RDONLY )
    printf( "read only" );
else if( accmode == O_WRONLY )
    printf( "write only" );
else if( accmode == O_RDWR )
    printf( "read write" );
else
{
    fprintf( stderr, "unkown access mode" );
    exit(1);
}

if( val & O_APPEND )
    printf( ", append" );
if( val & O_NONBLOCK )
    printf( ", nonblocking" );
if( val & O_SYNC )
    printf( ", synchronous writes" );
putchar( '\n' );
exit(0);
}
```

16

## Example #2: fcntl

```
#include <stdio.h>
#include <sys/types.h>
#include <fcntl.h>

/* flags are file status flags to turn on */
void set_fl( int fd, int flags )
{
    int val;

    if( (val = fcntl( fd, F_GETFL, 0 )) < 0 )
    {
        perror( "fcntl F_GETFL error" );
        exit( 1 );
    }

    val |= flags; /* turn on flags */
    if( fcntl( fd, F_SETFL, val ) < 0 )
    {
        perror( "fcntl F_SETFL error" );
        exit( 1 );
    }
}
```

17

## errno and perror()

- Unix provides a globally accesible integer variable that contains an error code number
- Error variable: *errno* – *errno.h*
- *perror*( " a string " ): a library routine

more /usr/include/asm/\*errno.h

```
#ifndef _I386_ERRNO_H
#define _I386_ERRNO_H

#define EPERM 1 /* Operation not permitted */
#define ENOENT 2 /* No such file or directory */
#define ESRCH 3 /* No such process */
#define EINTR 4 /* Interrupted system call */
#define EIO 5 /* I/O error */
#define ENXIO 6 /* No such device or address */
```

18

## errno and perror()

```
// file foo.c
#include <fcntl.h>
#include <unistd.h>
#include <stdio.h>

int main()
{
    extern int errno;
    int fd;

    /* open file "data" for reading */
    if( fd = open( "nosuchfile", O_RDONLY ) == -1 )
    {
        fprintf( stderr, "Error %d\n", errno );
        perror( "hello" );
    }
    /* end main */
}

(dkl:57) gcc foo.c -o foo
(dkl:58) ./foo
Error 2
hello: No such file or directory
```

19

## The Standard IO Library

- fopen, fclose, printf, fprintf, sprintf, scanf, fscanf,getc,putc,gets,fgets, etc.
- #include <stdio.h>

20

## Why use read()/write()

- Maximal performance
  - IF you know exactly what you are doing
  - No additional hidden overhead from stdio
- Control exactly what is written/read at what times

21

## File Concept – An Abstract Data Type

- File Types
- File Operations
- File Attributes
- File Structure - Logical
- Internal File Structure

22

## File Types

- Regular files
- Directory files
- Character special files
- Block special files
- FIFOs
- Sockets
- Symbolic Links

23

## File Operations

- Creating a file
- Writing a file
- Reading a file
- Repositioning within a file
- Deleting a file
- Truncating a file

24

## Files Attributes

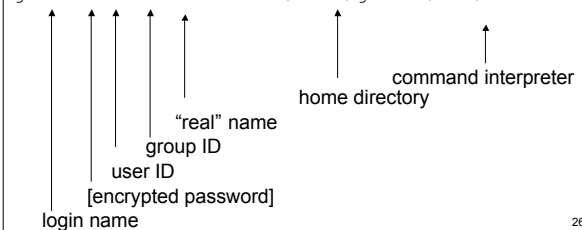
- Name
- Type
- Location
- Size
- Protection
- Time, date and user identification

25

## Users and Ownership: /etc/passwd

- Every File is owned by one of the system's users – identity is represented by the user-id (UID)
- Password file associate UID with system users.

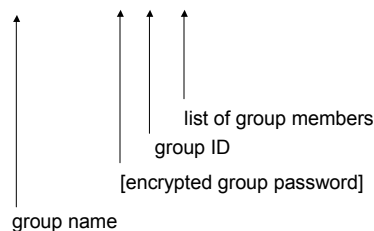
gates:x:65:20:B. Gates:/home/gates:/bin/ksh



26

## /etc/group

- Information about system groups
- faculty:x:23:maria,eileen,dkl



27

## Real uids

- The uid of the user who *started* the program is used as its *real uid*.
- The real uid affects what the program can do (e.g. create, delete files).
- For example, the uid of /usr/bin/vi is root:
 

```
- $ ls -alt /usr/bin/vi
lrwxrwxrwx 1 root root 20 Apr 13...
```
- But when I use `vi`, its *real uid* is `dkl` (not `root`), so I can only edit my files.

28

## Effective uids

- Programs can change to use the *effective uid*
  - the uid of the program *owner*
  - e.g. the `passwd` program changes to use its effective uid (`root`) so that it can edit the `/etc/passwd` file
- This feature is used by many system tools, such as logging programs.

29

## Real and Effective Group-ids

- There are also real and effective group-ids.
- Usually a program uses the *real group-id* (i.e. the *group-id of the user*).
- Sometimes useful to use *effective group-id* (i.e. group-id of program *owner*):
  - e.g. software shared across teams

30

## Extra File Permissions

- **Octal Value**      **Meaning**  
   04000      Set user-id on execution.  
              Symbolic: --s --- ---  
  
   02000      Set group-id on execution.  
              Symbolic: --- --s ---
- These specify that a program should use the effective user/group id during execution.
- For example:  
   -\$ ls -alt /usr/bin/passwd  
   -rwsr-xr-x 1 root root 25692 May 24.. 21

## Sticky Bit

- **Octal**      **Meaning**  
   01000      Save text image on execution.  
              Symbolic: --- --- --t
- This specifies that the program code should stay resident in memory after termination.  
   – this makes the start-up of the next execution faster
- Obsolete due to virtual memory.

32

## The superuser

- Most sys. admin. tasks can only be done by the *superuser* (also called the *root* user)
- Superuser
  - has access to all files/directories on the system
  - can override permissions
  - owner of most system files
- Shell command: `su <username>`
  - Set current user to superuser or another user with proper password access

33

## File Mode (Permission)

- S\_IRUSR -- user-read
- S\_IWUSR -- user-write
- S\_IXUSR -- user-execute
- S\_IRGRP -- group-read
- S\_IWGRP -- group-write
- S\_IXGRP -- group-execute
- S\_IROTH -- other-read
- S\_IWOTH -- other-write
- S\_IXOTH -- other-execute

34

## User Mask: *umask*

- Unix allows “masks” to be created to set permissions for “newly-created” directories and files.
- The `umask` command automatically sets the permissions when the user creates directories and files (`umask` stands for “user mask”).
- Prevents permissions from being accidentally turned on (hides permissions that are available).
- Set the bits of the `umask` to permissions you want to mask out of the file permissions.
- This process is useful, since user may sometimes forget to change the permissions of newly-created files or directories.

35

## `umask`: Calculations (1)

### • Defaults

File Type	Default Mode
Non-executable files	666
Executable files	777
Directories	777

From this initial mode, Unix “ands” the value of the `umask`.

36

## umask: Calculations (2)

- If you want a file permission of 644 (by default, without manually executing chmod) on a regular file, the umask would need to be 022.

Default Mode	666
umask	<u>-022</u>
New File Mode	644

- Bit level:  $\text{new\_mask} = \text{mode} \& \sim\text{umask}$

```
umask = 000010010 = ---rw-rw = 0022
~umask = 111101101
mode = 110110110 = rw-rw-rw = 0666
new_mask = 111100100 = rw----- = 0600
```

37

## umask

```
#include <sys/types.h>
#include <sys/stat.h>
mode_t umask( mode_t mask );
```

- Set file mode creation *mask* and return the old value.
- When creating a file, permissions are turned off if the corresponding bits in *mask* are set.
- Return value
  - This system call always succeeds and the previous value of the mask is returned.
  - cf. "umask" shell command

38

## Example: umask

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int main(void)
{
    umask(0);

    if( creat( "foo",
        S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH ) < 0 )
    {
        perror("creat error for foo");
        exit(1);
    }

    umask( S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH );
    if( creat( "bar",
        S_IRUSR|S_IWUSR|S_IRGRP|S_IWGRP|S_IROTH|S_IWOTH ) < 0 )
    {
        perror("creat error for bar");
        exit(1);
    }
    exit(0);
}
```

```
{saffron:maria:68} ls -lra foo bar
-rw-rw-rw- 1 dkl faculty  0 Apr 1 20:35 foo
-rw----- 1 dkl faculty  0 Apr 1 20:35 bar
```

39

## chmod and fchmod

```
#include <sys/types.h>
#include <sys/stat.h>
int chmod( const char *path, mode_t mode );
int fchmod( int fd, mode_t mode );
```

- Change permissions of a file.
- The mode of the file given by *path* or referenced by *fd* is changed.
- *mode* is specified by OR'ing the following.
  - $S_I\{R,W,X\}$  (USR,GRP,OTH) (basic permissions)
  - $S_ISUID, S_ISGID, S_ISVTX$  (special bits)
- Effective uid of the process must be zero (superuser) or must match the owner of the file.
- On success, zero is returned. On error, -1 is returned.

40

## Example: chmod

```
/* set absolute mode to "rw-r--r--" */
if( chmod("bar", S_IRUSR|S_IWUSR|S_IRGRP|S_IROTH) < 0 )
{
    perror("chmod error for bar");
    exit(1);
}
exit(0);
}
```

41

## chown, fchown, lchown

```
#include <sys/types.h>
#include <unistd.h>
int chown( const char *path, uid_t owner, gid_t group );
int fchown( int fd, uid_t owner, gid_t group );
int lchown( const char *path, uid_t owner, gid_t group );
```

- The owner of the file specified by *path* or by *fd*.
- Only the superuser may change the owner of a file.
- The owner of a file may change the group of the file to any group of which that owner is a member.
- When the owner or group of an executable file are changed by a non-superuser, the  $S_ISUID$  and  $S_ISGID$  mode bits are *cleared*.

42

## Obtaining File Information

- stat(), fstat(), lstat() For analyzing files.
- Retrieve all sorts of information about a file
  - Which device it is stored on
  - Don't need access right to the file, but need search rights to directories in path leading to file
  - Information:
    - Ownership/Permissions of that file,
    - Number of links
    - Size of the file
    - Date/Time of last modification and access
    - Ideal block size for I/O to this file

43

## struct stat

```
struct stat
{
    dev_t st_dev;      /* device num. */
    dev_t st_rdev;     /* device # spcl files */
    ino_t st_ino;      /* i-node num. */
    mode_t st_mode;     /* file type, mode, perms */
    nlink_t st_nlink;  /* num. of links */
    uid_t st_uid;      /* uid of owner */
    gid_t st_gid;      /* group-id of owner */
    off_t st_size;     /* size in bytes */
    time_t st_atime;   /* last access time */
    time_t st_mtime;   /* last mod. time */
    time_t st_ctime;   /* last stat chg time */
    long st_blksize;   /* best I/O block size */
    long st_blocks;    /* # of 512 blocks used */
}
```

We will look  
at st\_mode in detail.

44

## Recall: File Types

1. Regular File (text/binary)
2. Directory File
3. Character Special File  
e.g. I/O peripherals, such as /dev/tty0
4. Block Special File  
e.g. cdrom, such as /dev/mcd
5. FIFO (named pipes)
6. Sockets
7. Symbolic Links

45

## File Mix on a Typical System

<u>File Type</u>	<u>Count</u>	<u>Percentage</u>
regular file	30,369	91.7%
directory	1,901	5.7
symbolic link	416	1.3
char special	373	1.1
block special	61	0.2
socket	5	0.0
FIFO	1	0.0

46

## st\_mode Field

- This field contains type and permissions (12 lower bits) of file in bit format.
- It is extracted by AND-ing the value stored there with various constants
  - see man stat
  - also <sys/stat.h> and <linux/stat.h>
  - some data structures are in <bits/stat.h>

47

## Getting the Type Information

- AND the st\_mode field with S\_IFMT to get the type bits.
- Test the result against:
  - S\_IFREG Regular file
  - S\_IFDIR Directory
  - S\_IFSOCK Socket
  - etc.

48



## Example

```
struct stat sbuf;
:
if( stat( file, &sbuf ) == 0 )
    if( (sbuf.st_mode & S_IFMT) ==
        S_IFDIR )
        printf("A directory\n");
```

49

## Type Info. Macros

- Modern UNIX systems include test macros in `<sys/stat.h>` and `<linux/stat.h>`:

- `S_ISREG()` regular file
- `S_ISDIR()` directory file
- `S_ISCHR()` char. special file
- `S_ISBLK()` block special file
- `S_ISFIFO()` pipe or FIFO
- `S_ISLNK()` symbolic link
- `S_ISSOCK()` socket

50

## Example

```
struct stat sbuf;
:
if( stat(file, &sbuf ) == 0 )
{
    if( S_ISREG( sbuf.st_mode ) )
        printf( "A regular file\n" );
    else if( S_ISDIR(sbuf.st_mode) )
        printf( "A directory\n" );
    else ...
}
```

51

## Getting Mode Information

- AND the `st_mode` field with one of the following masks and test for non-zero:

- `S_ISUID` set-user-id bit is set
- `S_ISGID` set-group-id bit is set
- `S_ISVTX` sticky bit is set

- Example:

```
if( (sbuf.st_mode & S_ISUID) != 0 )
    printf("set-user-id bit is set\n");
```

52

## Getting Permission Info.

- AND the `st_mode` field with one of the following masks and test for non-zero:

- <code>S_IRUSR</code>	0400	user read
<code>S_IWUSR</code>	0200	user write
<code>S_IXUSR</code>	0100	user execute
- <code>S_IRGRP</code>	0040	group read
<code>S_IWGRP</code>	0020	group write
<code>S_IXGRP</code>	0010	group execute
- <code>S_IROTH</code>	0004	other read
<code>S_IWOTH</code>	0002	other write
<code>S_IXOTH</code>	0001	other execute

53

## Example

```
• struct stat sbuf;
:
printf( "Permissions: " );
if( (sbuf.st_mode & S_IRUSR) != 0 )
    printf( "user read, " );
if( (sbuf.st_mode & S_IWUSR) != 0 )
    printf( "user write, " );
:
```

54